
Application of the Lottery Ticket Hypothesis in NLP and Early Pruning (Proposal)

Anwendung der "Lottery Ticket"-Hypothese in NLP und frühem Pruning (Proposal)

Bachelor-Arbeit

Tim Unverzagt

KOM-type-number ???



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Fachbereich Informatik
Fachbereich ??? (Zweitmitglied)

Fachgebiet Natural Language
Processing
| Gutachter |

Application of the Lottery Ticket Hypothesis in NLP and Early Pruning (Proposal)
Anwendung der "Lottery Ticket"-Hypothese in NLP und frühem Pruning (Proposal)

Bachelor-Arbeit
Studiengang: Computational Engineering
KOM-type-number ???

Eingereicht von Tim Unverzagt
Tag der Einreichung: dd. month yyyy

Gutachter: ???
Betreuerin: Anna Filighera

Technische Universität Darmstadt
Fachbereich Informatik
Fachbereich ??? (Zweitmitglied)

Fachgebiet Natural Language Processing (KOM)
| Gutachter |

Erklärung zur Abschlussarbeit gemäß § 23 Abs. 7 APB der TU Darmstadt

Hiermit versichere ich, Tim Unverzagt, die vorliegende Bachelor-Arbeit ohne Hilfe Dritter und nur mit den angegebenen Quellen und Hilfsmitteln angefertigt zu haben. Alle Stellen, die Quellen entnommen wurden, sind als solche kenntlich gemacht worden. Diese Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen.

Mir ist bekannt, dass im Falle eines Plagiats (§38 Abs.2 APB) ein Täuschungsversuch vorliegt, der dazu führt, dass die Arbeit mit 5,0 bewertet und damit ein Prüfungsversuch verbraucht wird. Abschlussarbeiten dürfen nur einmal wiederholt werden.

Bei der abgegebenen Bachelor-Arbeit stimmen die schriftliche und die zur Archivierung eingereichte elektronische Fassung überein.

Darmstadt, den dd. month yyyy

Tim Unverzagt



Contents

1. Introduction	3
1.1. Motivation	3
1.2. Problem Statement and Contribution	3
1.3. Outline	3
2. Background	5
2.1. Basics of Neural Networks	5
2.2. Pruning	7
2.3. Basics of Natural Language Processing	7
3. Related Work	9
3.1. State of the art: Image Classification	9
3.2. Pruning	10
3.3. State of the art: Topic Classification	10
3.4. Early Pruning	10
3.5. Additions to the Lottery Ticket Hypothesis	11
4. Design	13
4.1. Requirements and Assumptions	13
4.2. Early Masking	13
4.2.1. Lenet-300-100	13
4.2.2. Conv-6	13
4.2.3. FGG-19	13
4.2.4. Early Stopping	14
4.3. Transfer to NLP	14
4.3.1. Based Module	14
4.3.2. Ensemble	14
4.4. Summary	14
5. Implementation	15
5.1. Design Decisions	15
5.2. Architecture	15
5.3. Interaction of Components	16
5.4. Summary	16
6. Data Sets	17
6.1. MNIST	17
6.2. CIFAR-10	17
6.3. 20-Newsgroup	17
6.4. Reuters-21578	17
7. Evaluation	19
7.1. Goal and Methodology	19
7.2. Evaluation Setup	19
7.3. Evaluation Results	19

7.4. Analysis of Results	19
8. Conclusions	21
8.1. Summary	21
8.2. Contributions	21
8.3. Future Work	21
8.4. Final Remarks	21
Bibliography	22
Appendices	25
A. A history of neural networks	27

Abstract

The abstract goes here...



1 Introduction

1.1 Motivation

- LTH has demonstrated extreme pruning on different architectures
- Study of lottery-ticket emergence points might result in a reasoned early pruning approach
- LTH might bring a new and promising pruning approach to NLP

1.2 Problem Statement and Contribution

- Calculate pruning masks earlier in training and check if LTH still holds.
Observing when lottery-tickets are no longer found might improve understanding of early pruning.
- Research whether the point of lottery-ticket-emergence can be estimated "a-priori".
- Implement an architecture comparable to the ones studied in the lottery-ticket hypothesis performing well on an NLP-task with similar structure.
- Determine whether the LTH holds on said architecture.

1.3 Outline

???



2 Background

2.1 Basics of Neural Networks

Neural networks are a part of most major AI-breakthrough in the last decade enabling computers to compete in fields formerly championed by humans.¹ They implement a statistical understanding of AI, which is to say that they try to find a specific model optimizing the likelihood of reproducing input-output pairs similar to some training data. The competing philosophy directly divines behaviour rules, frequently from expert knowledge, and as such is far less dependant from data. [citation needed] For the former concept its model classes are the essential point of design. A multitude of properties maybe sought after in a model class of which a few important ones are:

- **Richness:**

The diversity of single models in the class and thus the ability to fit a wide field of different input-output landscapes.²

If a model class is inherently restricted the underlying relation between inputs and outputs might simply be beyond the expressive capabilities of all its models.

In other words: If a model class is not rich enough all of its models will underfit the given training data.

- **Stability:**

Tendency of similar models in the class to handle inputs in a similar way.

If your model class shows unstable behavior defining a sensible way to search it for good models becomes difficult.

- **Interpretability of Models:**

Ease of formulating knowledge out of any given model in the class.

As fields exist in which statistical AI outperform experts the extraction of knowledge understandable and applicable by humans is of special interest.

- [citation needed]

If one knows an entity that already performs well on a given task it is a sensible approach to design ones model class to reproduce its decision process. Humans usually are such entities for many tasks of interest to AI research so they are a natural source of inspiration. Neural networks essentially are simplified models of a human central nervous system.

The most basic building block of the human central nervous system is a neuron which can receive multiple stimuli and is able to produce an output if the combined stimulation exceeds a threshold. [citation needed] One such neuron and its stimulus measure are depicted in 2.1. Another functionality observed in nature is the ability of a neuron to strengthen the connection to any source of stimulus thus giving said source more influence on whether the neuron produces an output. [citation needed]

The canonical mathematical model of a neuron, as seen in 2.2, is defined as:

1

- 2011: "Watson" of IBM defeats two former grand champions in "Jeopardy!" [LF11]
- 2011: "Siri" enables users to use natural language to interact with their phones [Aro11]
- 2015: A convolutional neural network classifies images from the ImageNet dataset more accurately than human experts [RDS⁺15] [HZRS15]
- 2016: "AlphaGo" beats Lee Sedol, one of the world's strongest Go players [Gib16] [SSS⁺17]

² More formally the richness of a model class can be described as the amount of different functions from the input-space to the output-space which can be expressed through a model of said class.




Figure 2.1.: Representation of a biological Neuron [CDC18] edited




Figure 2.2.: Abstraction of a Neuron [DMK⁺12]




Figure 2.3.: A small fully-connected network [Bel18]

- **Inputs x_i :**
All stimuli of a neuron are simply referred to as its inputs
- **Weights w_i :**
The ability to assign importances is modelled as weights which are coupled to specific stimuli
- **Combined Weighted Inputs $\sum_{i=1}^n w_i x_i$:**
After the inputs are scaled by their according weight they superpose to form the total excitation of the neuron
- **Activation Function $\Phi(\sum_{i=1}^n w_i x_i)$:**
...
- **Bias b :**
...

As an individual neurons is too simple to model any complex relations between inputs and outputs the next step is to aggregate multiple neurons. Figure 2.3 displays a few neurons coming together to form a simple fully-connected feed-forward network.³

- FNN as universal approximator (but overfitting)
- Issue of computational expense
- CNNs

³ Inputs of neural networks are often called "features" and fully-connected networks are frequently referred to as "dense"

2.2 Pruning

- Generalization (Anti-Overfitting)
- weight magnitude based pruning
- Levels of pruning

2.3 Basics of Natural Language Processing

- Preprocessing (Padding)
- Quantifying text (Features, ngrams, BoW)
- Language models



3 Related Work

To quantify the goals previously defined the context of current research is needed. The importance of any work assuming an underlying architecture can not be correctly evaluated without knowledge about the quality of said architecture. As such this section shortly presents state-of-the-art approaches to the tasks relevant to this thesis. Additionally an overview over previous compression methods and their achievements is given.

3.1 State of the art: Image Classification

MNIST and CIFAR-10 are both datasets containing small images which are to be classified according to the object they display. While MNIST contains gray scale images of hand-written digits CIFAR-10 consists of colorful real-world images. State of the art approaches deliver nigh perfect superhuman accuracy on both data sets.

For MNIST Kowasari et al. with their random multi-purpose deep learning ensemble report the nominal highest performance [KHB⁺18] although many others achieve similar results through varying means. Already in 2012 Ciresan et al. describe a deep and sparse convolutional architecture that resembles the visual cortex of mammal [CMS12]. Later Sato et al. apply data-augmentation [SNY15], Chang Jia-Ren & Chen Yong-Sheng package whole architectures and treat them like layers [CC15] and Hasanpour et al. carefully design a small and simple convolutional network through the use of structural heuristics [HRFS16] all reproducing the same performance.

In contrast the three best-performing approaches to CIFAR-10 are all published in 2019. Currently an ensemble of auto-encoding transformations claims the highest performance. Wang et al. provide their model with a rich class of transformations to help it abstract the input. [WKLQ19]. Close second and third are Cai et al. with a direct network-architecture-search scheme [CZH18] and Hu et al. with a novel network building block that explicitly models interaction between channels [HSS17].

While Frankle & Carbin do not provide exact values in the LTH-paper their figures indicate that they reproduce roughly 98% accuracy on MNIST and 90% on CIFAR-10 [FC18].¹ This result is reproducible with the source code provided alongside this thesis.

Accuracy %	MNIST	CIFAR-10	Published
EnAET		98.0	2019
DirNAS		97.9	2019
Squee		97.88	2019
RMDL	99.82	91.2	2018
Simple	99.8	95.5	2016
BatchNorm	99.8	93.3	2015
APAC	99.8	89.7	2015
Multi-Column	99.8	88.8	2012
Lenet-FCN	~98		LTH
VGG-19		~90	LTH

Figure 3.1.: Performance for Image Classification

¹ State-of-the-Art architectures are presented only if no extra training data was used and as described on <https://paperswithcode.com/sota>

Accuracy %	20-News	Reuters	Published
Neural BoE	88.1		2019
Graph Star	86.9		2019
RMDL		90.69	2018
multi-scale CNN	86.12		2018

Figure 3.2.: Performance for Topic Classification

3.2 Pruning

Beginning around 1990 with M.C. Mozer & P. Smolensky [MS89] as well as LeCun et al. [LDS90] weights were being removed from neural networks after training them for a task. Shortly thereafter the idea of further training a pruned network was proposed [HS93] which became common practice over the next decade. While LeCun et al. describe a network compression factor of $\times 4$, more recent works achieve a factor of $\times 9$ to $\times 16.6$ while loosing no or close to no accuracy [HPTD15] [LWL17].

In their paper on the Lottery-Ticket-Hypothesis (now LTH) J. Frankle & M. Carbin report pruning over 98,5% of weights in one of their networks while maintaining network capabilities which amounts to a compression rate of over $\times 50$

3.3 State of the art: Topic Classification

In the field of NLP topic classification is arguably the task most similar to image classification and Reuters-21578 is arguably the most iconic dataset for such a task. Yet neither do its corresponding state of the art architectures compare sensibly to the ones studied by Frankle & Carbin nor is Reuters-21578 structurally akin to MNIST. The essential differences will be covered in section 6.

20-Newsgroup is another NLP data set which not only aligns more with MNIST and CIFAR-10 but for which an CNN architecture exists. In their work Pappagari et al. develop an approach integrating the implicit verification objective and applying multiple language models to the different channels of a CNN [PVD18]. They come close to state of the art performance on 20-Newsgroup.

3.4 Early Pruning

In a recent paper [LSZ⁺18] Z.Liu et al. observe that if pruned networks are trained with randomly reinitialized weights instead of fine-tuning their previous ones they retain from the original network, the pruned networks keep their capabilities. They conclude that said weights can not be essential to a pruned networks quality, contrary to prior common belief. Thus Z.Liu et al. claim that the architecture of pruned networks is responsible for its capabilities and furthermore that pruning can be interpreted as a kind of network architecture search .

After the effectiveness of pruning is established and its interpretation as network architecture search becomes available there is a legitimate question whether all the weights in a network are really necessary for all of the training. In a paper of Y. Li & W. Zhao & L. Schang from early 2019 [LZS19] they describe a method named IPLT to prune common convolutional network architectures at the filter level and especially before convergence. Thus they do not only compress the networks by a factor of $\times 10$ but also speed up training by a similar magnitude. If the LTH can be applied in such a fashion a speed-up of up to $\times 20$ should be expected.

3.5 Additions to the Lottery Ticket Hypothesis

Even though the Lottery-Ticket-Hypothesis was only proposed earlier this year additional papers on the topic exist. In a paper from June 2019 J. Frankle & M. Carbin et al. [FDRC19] expand their method to find winning tickets on deep convolutional network architectures that proved difficult before. They attribute this achievement to the decision of not returning to the very first state of the network but to one a few iterations into training. Not only does this mark a lower limit for how early pruning is possible with the LTH but it also implies that a certain structure emerges after little training of the big network. Whether said structure only marks a point for valid reinitialization or rather already one for magnitude-based pruning is part of what this thesis wants to explore.

Additionally H. Zhou et al. [ZLLY19] document an ablation study on the phenomenon of lottery tickets. They reaffirm the initially naive magnitude-based pruning and describe "supermasks" that improve accuracy when applied to the initial network even without additional training. Finally they find that a replacement of all weights in the pruned network by a constant with same sign as said weights does not significantly influence the networks capabilities. As such H. Zhou et al. conclude that the sign of weights are the essential property for such neural networks.



4 Design

Hint:

This chapter should describe the design of the own approach on a conceptional level without mentioning the implementation details. The section should have a length of about five pages.

4.1 Requirements and Assumptions

???

4.2 Early Masking

4.2.1 Lenet-300-100

- FC: 300, 100, 10
- Prune 20%
- Optimizer: Adam $3 \cdot 10^{-4}$

4.2.2 Conv-6

- Con: 64, 64, pool
- Con: 128, 128, pool
- Con: 256, 256, pool
- FC: 256, 256, 10
- Prune 20% FC
- Prune 15% Con
- Optimizer: Adam $3 \cdot 10^{-4}$

4.2.3 FGG-19

All convolutions are 3x3 with 1 padding to preserve image-size. Pooling is 2x2-subsampling with stride 2. With each pooling the channel-depth is doubled

- Con: 64, 64, pool
- Con: 128, 128, pool
- Con: 256, 256, 256, 256, pool
- Con: 512, 512, 512, 512, pool
- Con: 512, 512, 512, 512, avg-pool
- FC: 10
- Prune 20% Con
- Optimizer: Momentum 0.9

4.2.4 Early Stopping

A-posteriori vs. A-priori

4.3 Transfer to NLP

4.3.1 Based Module

- wordEmbed: at different scales (?) 300 dims
- Con: (?) 3 filter maps (width: 1:3:22)
- TempPool: 2 | 7
- Dropout
- GlobalAvePool

4.3.2 Ensemble

- Module: 16x
- Concat:
- Dropout
- Output
- Loss: binary cross-entropy

4.4 Summary

5 Implementation

Hint:

This chapter should describe the details of the implementation addressing the following questions:

1. What are the design decisions made?
2. What is the environment the approach is developed in?
3. How are components mapped to classes of the source code?
4. How do the components interact with each other?
5. What are limitations of the implementation?

The section should have a length of about five pages.

5.1 Design Decisions

- tensorflow 2.0
- scikit-learn
- pickle

5.2 Architecture

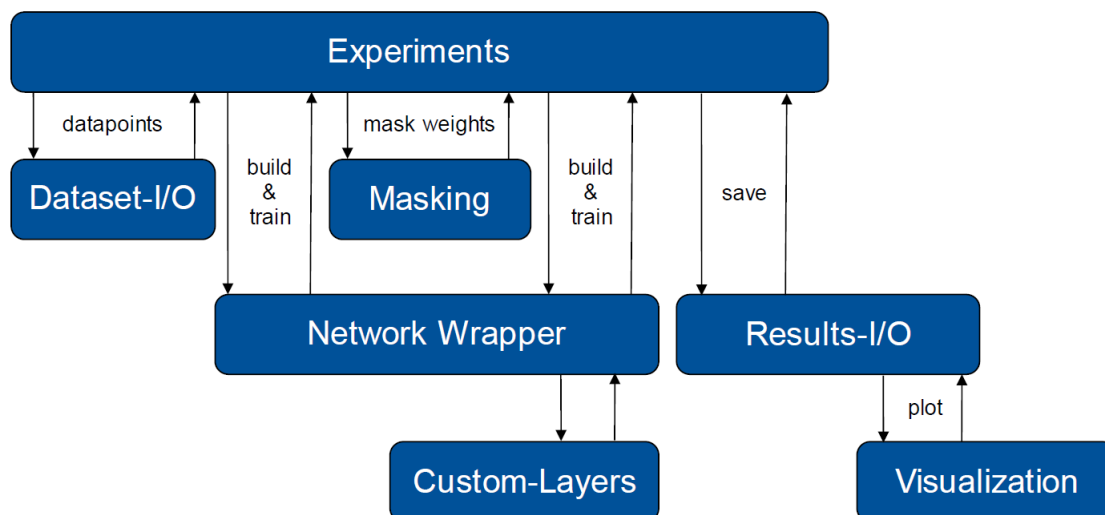


Figure 5.1.: project architecture

5.3 Interaction of Components




Figure 5.2.: project architecture

5.4 Summary

6 Data Sets

6.1 MNIST

The MNIST-dataset contains 25x25 gray-scale images of handwritten digits padded to 28x28 [YL].

6.2 CIFAR-10

6.3 20-Newsgroup

6.4 Reuters-21578

The Reuters-21578-dataset contains 21578 articles published by the Reuters News Agency in 1987 [Lew]. Reuters-21578 differs from the previous data sets in the sense that it lacks a few fundamental properties. In particular Reuters-21578 is not only multi-class but rather multi-label meaning that any one data point can satisfy multiple categories. Additionally there are categories in Reuters-21578 that have no associated positive example and even for all remaining ones the amount of samples is heavily skewed. In order to restore parts of the missing properties with minimal change to the dataset different subsets of Reuters-21578 have been chosen by different researchers.

F. Debole & F. Sebastiani [DS05] describe those subsets, starting out stating that close to half of the data points are unusable which leaves 12,902 documents. 9,603 are marked for training and 3,299 for validation.¹ They also point out the different groups of categories used for classification:

- **R(115)**
The group with the 115 categories containing at least one positive training example.
- **R(90)**
The group with the 90 categories containing at least one positive training and test example.
- **R(10)**
The group with the 10 categories containing the most examples.

	MNIST	CIFAR-10	20-Newsgroup	Reuters-21578
N. labels	10	10	20	10 to 115
N. datapoints	70.000	60.000	18846	12.902
fixed split	x	x	"bydate"	"ModApté"
shortened			x	x
class imbalance				x
multi-label				x

¹ While different training-splits were used for Reuters-21578 "ModApté" has become the canonical choice



7 Evaluation

Hint:

This chapter should describe how the evaluation of the implemented mechanism was done.

1. Which evaluation method is used and why? Simulations, prototype?
2. What is the goal of the evaluation? Comparison? Proof of concept?
3. Which metrics are used for characterizing the performance, costs, fairness, and efficiency of the system?
4. What are the parameter settings used in the evaluation and why? If possible always justify why a certain threshold has been chosen for a particular parameter.
5. What is the outcome of the evaluation?

The section should have a length of about five to ten pages.

7.1 Goal and Methodology

- Early Pruning:
Proof of concept
+ trial of select early stopping points
- Transfer to NLP:
Proof of concept

7.2 Evaluation Setup

- Early Pruning:
LTH holds for mask with less than 50% of iterations to early stop
- Transfer to NLP:
Network holds performance to within 1%-point while at least 50% weights are pruned

7.3 Evaluation Results

7.4 Analysis of Results



8 Conclusions

Hint:

This chapter should summarize the thesis and describe the main contributions of the thesis. Subsequently, it should describe possible future work in the context of the thesis. What are limitations of the developed solutions? Which things can be improved? The section should have a length of about three pages.

8.1 Summary

8.2 Contributions

8.3 Future Work

8.4 Final Remarks



Bibliography

- [Aro11] Jacob Aron. How innovative is apple’s new voice assistant, siri? *New Scientist*, 212(2836):24, 2011.
- [Bel18] Soufiane Belharbi. Neural networks regularization through representation learning, 07 2018.
- [CC15] Jia-Ren Chang and Yong-Sheng Chen. Batch-normalized maxout network in network. *CoRR*, abs/1511.02583, 2015.
- [CDC18] Mary Ann Clark, Matthew Douglas, and Jung Choi. *Biology 2e*. OpenStax, 2018.
- [CMS12] Dan C. Ciresan, Ueli Meier, and Jürgen Schmidhuber. Multi-column deep neural networks for image classification. *CoRR*, abs/1202.2745, 2012.
- [CZH18] Han Cai, Ligeng Zhu, and Song Han. Proxylessnas: Direct neural architecture search on target task and hardware. *CoRR*, abs/1812.00332, 2018.
- [DMK⁺12] Jelena Djuris, Djordje Medarević, Marko Krstić, Ivana Vasiljević, Ivana Aleksić, and Svetlana Ibrić. Design space approach in optimization of fluid bed granulation and tablets compression process. *TheScientificWorldJournal*, 2012:185085, 07 2012.
- [DS05] Franca Debole and Fabrizio Sebastiani. An analysis of the relative hardness of reuters-21578 subsets. *Journal of the American Society for Information Science and Technology*, 56(6):584–596, 2005.
- [FC18] Jonathan Frankle and Michael Carbin. The lottery ticket hypothesis: Training pruned neural networks. *CoRR*, abs/1803.03635, 2018.
- [FDRC19] Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M. Roy, and Michael Carbin. The lottery ticket hypothesis at scale. *CoRR*, abs/1903.01611, 2019.
- [Gib16] Elizabeth Gibney. Google ai algorithm masters ancient game of go. *Nature News*, 529(7587):445, 2016.
- [HPTD15] Song Han, Jeff Pool, John Tran, and William Dally. Learning both weights and connections for efficient neural network. In C. Cortes, N. D. Lawrence, D. D. Lee, M. Sugiyama, and R. Garnett, editors, *Advances in Neural Information Processing Systems 28*, pages 1135–1143. Curran Associates, Inc., 2015.
- [HRFS16] Seyyed Hossein HasanPour, Mohammad Rouhani, Mohsen Fayyaz, and Mohammad Sabokrou. Lets keep it simple, using simple architectures to outperform deeper and more complex architectures. *CoRR*, abs/1608.06037, 2016.
- [HS93] Babak Hassibi and David G Stork. Second order derivatives for network pruning: Optimal brain surgeon. In *Advances in neural information processing systems*, pages 164–171, 1993.
- [HSS17] Jie Hu, Li Shen, and Gang Sun. Squeeze-and-excitation networks. *CoRR*, abs/1709.01507, 2017.
- [HZRS15] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. In *The IEEE International Conference on Computer Vision (ICCV)*, December 2015.

-
- [KHB⁺18] Kamran Kowsari, Mojtaba Heidarysafa, Donald E. Brown, Kiana Jafari Meimandi, and Laura E. Barnes. Rmdl: Random multimodel deep learning for classification. In *Proceedings of the 2Nd International Conference on Information System and Data Mining, ICISDM '18*, pages 19–28, New York, NY, USA, 2018. ACM.
- [Kim14] Yoon Kim. Convolutional neural networks for sentence classification. *CoRR*, abs/1408.5882, 2014.
- [LDS90] Yann LeCun, John S Denker, and Sara A Solla. Optimal brain damage. In *Advances in neural information processing systems*, pages 598–605, 1990.
- [Lew] David D. Lewis. Reuters-21578. <http://www.daviddlewis.com/resources/testcollections/reuters21578/>.
- [LF11] Adam Lally and Paul Fodor. Natural language processing with prolog in the ibm watson system. *The Association for Logic Programming (ALP) Newsletter*, 2011.
- [LSZ⁺18] Zhuang Liu, Mingjie Sun, Tinghui Zhou, Gao Huang, and Trevor Darrell. Rethinking the value of network pruning. *CoRR*, abs/1810.05270, 2018.
- [LWL17] Jian-Hao Luo, Jianxin Wu, and Weiyao Lin. Thinet: A filter level pruning method for deep neural network compression. In *The IEEE International Conference on Computer Vision (ICCV)*, Oct 2017.
- [LZS19] Yue Li, Weibin Zhao, and Lin Shang. Really should we pruning after model be totally trained? pruning based on a small amount of training. *CoRR*, abs/1901.08455, 2019.
- [MS89] Michael C Mozer and Paul Smolensky. Skeletonization: A technique for trimming the fat from a network via relevance assessment. In *Advances in neural information processing systems*, pages 107–115, 1989.
- [PVD18] R. Pappagari, J. Villalba, and N. Dehak. Joint verification-identification in end-to-end multi-scale cnn framework for topic identification. In *2018 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 6199–6203, April 2018.
- [RDS⁺15] Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg, and Li Fei-Fei. Imagenet large scale visual recognition challenge. *International Journal of Computer Vision*, 115(3):211–252, Dec 2015.
- [SNY15] Ikuro Sato, Hiroki Nishimura, and Kensuke Yokoi. APAC: augmented pattern classification with neural networks. *CoRR*, abs/1505.03229, 2015.
- [SSS⁺17] David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, et al. Mastering the game of go without human knowledge. *Nature*, 550(7676):354, 2017.
- [WKLQ19] Xiao Wang, Daisuke Kihara, Jiebo Luo, and Guo-Jun Qi. Enaet: Self-trained ensemble autoencoding transformations for semi-supervised learning, 2019.
- [YL] Christopher J.C.Burges Yann LeCun, Corinna Cortes. The mnist database. <http://yann.lecun.com/exdb/mnist/>.
- [ZLLY19] Hattie Zhou, Janice Lan, Rosanne Liu, and Jason Yosinski. Deconstructing lottery tickets: Zeros, signs, and the supermask. *CoRR*, abs/1905.01067, 2019.

Appendices



A A history of neural networks

- 1. wave: 1955-1970
- 2. wave: 1985-2000 ?
- 3. wave: ???




Figure A.1.: Relative amount of occurrences of the word "Perceptron" in published books between 1940 and 2009




Figure A.2.: Relative amount of occurrences of the word "Backpropagation" in published books between 1940 and 2009